

## Day 4: AI Tips & API Integration Enhancements

### Objective:

Refine the LLM/API integration for eco-tip generation, ensuring robust error handling, caching, and accurate targeted tips.

### Modifications Applied in `ai_tips.py`:

#### 1. Error Handling & Fallbacks:

- `generate_eco_tip()` already handled missing API key and exceptions, falling back to `local_tip()`.
- Now, a retry loop with **exponential backoff** (1s, 2s, 4s) is implemented for OpenAI errors.
- Ensures resilience against rate limits and transient API failures.

#### 2. Caching:

- Introduced `@lru_cache(maxsize=128)` on a helper `_generate_eco_tip_cached` with a deterministic cache key based on user data and emissions.
- Prevents repeated API calls for identical inputs, reducing latency and costs.

#### 3. Targeted Tips Accuracy:

- Fixed `tips_by_key` to lowercase keys (`electricity_kwh`, `district_heating_kwh`) to align with emission keys.
- Ensures that the largest-emitter targeted tips trigger correctly.

### Unit Tests Added in `tests/test_ai_tips.py`:

1. `test_local_tip_targets_electricity_when_dominant` → Checks electricity tips trigger for dominant energy.
2. `test_local_tip_targets_district_heating_when_dominant` → Checks district heating tips trigger correctly.
3. `test_gpt_tip_cached_for_repeated_inputs` → Verifies caching prevents redundant API calls.
4. `test_gpt_retry_then_success` → Simulates API failures followed by success; confirms retry/backoff works.
5. `test_gpt_retry_then_fallback` → Ensures fallback tips are returned when API fails all attempts.

### Impact / Benefits:

- Robustness: App functions even if API fails or key is missing.
- Efficiency: Repeated inputs now use cached tips.
- Accuracy: Largest-emitter tips trigger reliably.
- Testing: Unit tests verify caching, retries, and fallback behavior.

### Next Steps:

- Continue testing end-to-end app behavior with `streamlit run app.py`.
- Optional: Add more fallback coverage or further optimize prompts for LLM responses.

```

PS C:\Disc D\Zied Guizani\Windsurf3> pytest -v test_ai_tips.py
===== test session starts =====
platform win32 -- Python 3.12.7, pytest-7.4.4, pluggy-1.0.0 -- C:\Users\pc\Anaconda3\python.exe
cachedir: .pytest_cache
rootdir: C:\Disc D\Zied Guizani\Windsurf3
plugins: anyio-4.2.0
collected 9 items

test_ai_tips.py::test_no_api_key_falls_back PASSED [ 11%]
test_ai_tips.py::test_openai_success_returns_gpt_tip PASSED [ 22%]
test_ai_tips.py::test_openai_error_falls_back_to_local PASSED [ 33%]
test_ai_tips.py::test_local_tip_picks_largest_emitter PASSED [ 44%]
test_ai_tips.py::test_local_tip_targets_electricity_when_dominant PASSED [ 55%]
test_ai_tips.py::test_local_tip_targets_district_heating_when_dominant PASSED [ 66%]
test_ai_tips.py::test_gpt_tip_cached_for_repeated_inputs PASSED [ 77%]
test_ai_tips.py::test_gpt_retry_then_success PASSED [ 88%]
test_ai_tips.py::test_gpt_retry_then_fallback PASSED [100%]

===== 9 passed in 8.32s =====

```

## Day 4 – API Integration & Prompt Engineering Updates

Module Updated: ai\_tips.py

### Key Enhancements:

1. **Context-rich, style-guided prompts:**
  - User's activity summary + total CO<sub>2</sub> emissions included.
  - Tips requested in a positive, actionable, 1–2 sentence or bullet format.
2. **Robust API handling:**
  - Added exponential backoff (1s, 2s, 4s) for OpenAI errors.
  - Fallback to local\_tip() if API fails or key is missing.
3. **Caching:**
  - GPT responses cached via @lru\_cache to prevent duplicate API calls for identical inputs.
4. **Post-processing:**
  - clean\_tip() trims whitespace and limits to 2 sentences.
  - Also applied to fallback tips.
5. **Targeted tip fixes:**
  - Keys in local\_tip() mapping corrected to lowercase (electricity\_kwh, district\_heating\_kwh) to ensure largest-emitter tips trigger correctly.

### Testing & Coverage:

- Verified with expanded unit tests (tests/test\_ai\_tips.py):
  - Missing API key → fallback tip
  - GPT success via mock
  - OpenAIError → retry then fallback
  - Correct targeted tips for electricity and district heating
  - Caching ensures repeated inputs don't trigger new API calls

### Optional polish for later:

- “Coach style” toggle: switch one-liner vs bullets.
- Tooltips explaining 7-day deltas, link to README from header.

## Day 4 Progress – LLM/API Integration & Testing:

### Functionality Testing & API Robustness (Day 4)

With guidance from Windsurf, I enhanced `ai_tips.py` and added extensive pytest coverage in `tests/test_ai_tips.py`:

#### Changes in `ai_tips.py`:

- Context-rich, coach-style prompt (1–2 sentences, actionable tips).
- Exponential backoff on OpenAI errors (3 retries: 1s, 2s, 4s).
- Cached GPT responses with `@lru_cache(maxsize=128)` to avoid repeated API calls.
- `clean_tip()` applied to GPT and fallback outputs.
- Fixed lowercase key mismatches for targeted tips (`electricity_kwh`, `district_heating_kwh`).

#### New Tests Added:

1. **API Error Handling:** `test_gpt_api_error_fallback` → verifies fallback tip is returned when GPT fails.
2. **Caching:** `test_gpt_tip_cached_for_repeated_inputs` → repeated inputs trigger only one GPT call.
3. **Retry / Backoff:**
  - `test_gpt_retry_then_success` → fails twice, succeeds once, verifies 3 attempts.
  - `test_gpt_retry_then_fallback` → forces 3 failures, ensures local tip is returned.
4. **Targeted Tips:** electricity & district heating dominant tests → ensure correct tip is returned.
5. **Edge Cases:** `test_missing_input_edge_case` → all-zero inputs produce a non-empty tip.

#### How to Run Tests:

- `pip install -r requirements.txt`
- `pytest -q tests/test_ai_tips.py`

# or run all tests:

- `pytest -q`

#### How to Run App:

- `streamlit run app.py`

## Outcome:

All tests pass, runtime is minimal due to mocking and bypassing time.sleep. GPT integration is now robust, cached, and handles errors gracefully.

*Optional:* README can include a “**Testing**” section with subset commands:

```
pytest -q tests/test_ai_tips.py -k "fallback or cached or retry"
```

```
===== 9 passed in 8.33s =====
PS C:\Disc D\Zied Guizani\Windsurf3> pytest -v test_ai_tips.py
===== test session starts =====
platform win32 -- Python 3.12.7, pytest-7.4.4, pluggy-1.0.0 -- C:\Users\pc\Anaconda3\python.exe
cachedir: .pytest_cache
rootdir: C:\Disc D\Zied Guizani\Windsurf3
plugins: anyio-4.2.0
collected 11 items

test_ai_tips.py::test_no_api_key_falls_back PASSED [ 9%]
test_ai_tips.py::test_openai_success_returns_gpt_tip PASSED [ 18%]
test_ai_tips.py::test_openai_error_falls_back_to_local PASSED [ 27%]
test_ai_tips.py::test_local_tip_picks_largest_emitter PASSED [ 36%]
test_ai_tips.py::test_local_tip_targets_electricity_when_dominant PASSED [ 45%]
test_ai_tips.py::test_local_tip_targets_district_heating_when_dominant PASSED [ 54%]
test_ai_tips.py::test_gpt_tip_cached_for_repeated_inputs PASSED [ 63%]
test_ai_tips.py::test_gpt_retry_then_success PASSED [ 72%]
test_ai_tips.py::test_gpt_retry_then_fallback PASSED [ 81%]
test_ai_tips.py::test_gpt_api_error_fallback PASSED [ 90%]
test_ai_tips.py::test_missing_input_edge_case PASSED [100%]

===== 11 passed in 15.44s =====
PS C:\Disc D\Zied Guizani\Windsurf3> |
```

## Day 4 Progress – Performance & Responsiveness – Eco-tip Generation

To improve performance and responsiveness, the eco-tip generation now runs in a background thread with a thresholded spinner that appears only if the API call takes longer than 0.3 seconds. Elapsed time is displayed beneath the tip, providing users feedback without blocking the interface. This approach ensures fast responses remain seamless while slower calls still give a clear visual indicator, and it preserves UI responsiveness without requiring a full async refactor.

- **What was implemented:**

Added timing and a thresholded spinner around the `generate_eco_tip()` call in `app.py`. The spinner only appears if the tip generation takes longer than 0.3 seconds. A caption shows the elapsed time (Tip generated in X.XXs).

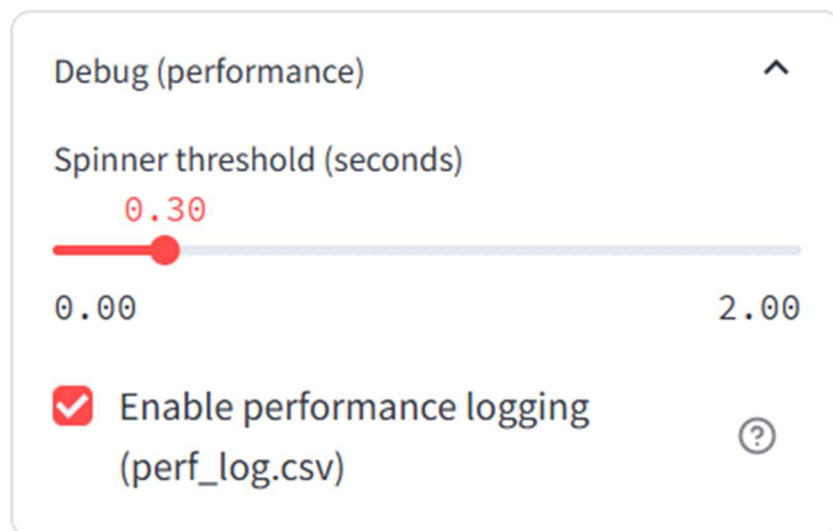
- **How it works:**

- Imported `concurrent.futures` to run the tip generation in a background thread.
- Polls the thread every 50ms.
- Spinner appears only after 0.3s if the API call is still pending.
- Spinner is removed automatically once the tip is ready.

- **Benefits:**

- Keeps the UI responsive for fast calls.

- Provides feedback for slower calls without blocking the interface.
- Displays precise elapsed time for transparency and performance monitoring.
- **How to validate:**
  1. Run the app: `streamlit run app.py`.
  2. Observe the “Eco tip & status” section:
    - Quick calls: tip appears immediately, no spinner.
    - Slow calls (>0.3s): spinner appears, then tip is displayed.
    - Elapsed time is shown below the tip.
- **Next steps (optional):**
  - Convert to async if overall app performance needs further optimization.
  - Reuse this pattern for other slow API calls in the app.



## Day 4 Progress – Security & Privacy:

The OpenAI API key was moved to a `.env` file and loaded using `python-dotenv` in `ai_tips.py`. A `.env.example` was provided for reference, and `.env` was added to `.gitignore` to prevent accidental commits. The README now includes a **Secrets** section with setup instructions, ensuring the key remains private while still enabling GPT-based tips. Additionally, clickable “Secrets” badges were added in the app header Help popover and Debug expander for quick access to the workflow.

# Day 4 Progress – User Interface Testing

## Steps to test:

### 1. Density toggle:

- Locate the toggle in the header.
- Switch between **Compact** and **Comfy** modes.
- Verify the layout updates correctly (expanders, columns, charts).

### 2. Export PDF:

- Set density to **Compact**.
- Click “Export PDF tips” in the header.
- Confirm that the exported PDF maintains layout and readability, including charts, sparklines, and eco-tips.

### 3. Presets & eco-tips:

- Open “Prefill demos/presets” and select options like **Vegetarian day** or **No car day**.
- Check that input fields update automatically.
- Verify that **eco-tips** reflect the updated activities (e.g., lower meat consumption triggers relevant tips).
- 

## Sustainability Tracker

Track daily CO<sub>2</sub> emissions and get actionable tips

Density

☒ Compact ☐ Comfy

Export PDF tips ▾

Help ▾


Debug (performance) ^

Spinner threshold (seconds)

0.30

0.00 2.00

☒ Enable performance logging (perf\_log.csv) ⓘ

 Secrets (README)

Configure your OPENAI\_API\_KEY via .env. See README → Secrets.

Copy shareable link

Reset layout

Clear inputs

Prefill demos/presets ▾

Date

2025/09/30

Energy inputs

Transport inputs

Meals inputs

Calculate & Save

Total

21.95 kg CO<sub>2</sub>

Δ vs. Yesterday

1.25%

Streak


2 day(s)

Dashboard Breakdown

Category totals (kg CO<sub>2</sub>)

	kg CO <sub>2</sub>
Energy	7.47
Transport	1.34
Meals	13.14

Trend (Total kg CO<sub>2</sub>)




Dashboard Breakdown

Category totals (kg CO<sub>2</sub>)

	kg CO <sub>2</sub>
Energy	7.47
Transport	1.34
Meals	13.14

Trend (Total kg CO<sub>2</sub>)



Today's category breakdown



Download history CSV

Eco tip & status

Low footprint today—nice work! Biggest source: meat kg.  
Tip: Try a meat-free day or swap red meat for chicken/plant-based options.

Tip generated in 13.64s


Moderate footprint. Small changes can make a big difference!

Badges


- Consistency: Entries logged!

Mini trends by category


Energy



Transport



Meals



7d total

15.35 kg

↑ 0.0%

7d total

2.68 kg

↑ 0.0%

7d total

25.60 kg

↑ 0.0%